

A woman with long dark hair, wearing a grey coat, stands in a city at night, looking down at a tablet she is holding. The background is filled with blurred lights from street lamps and buildings, creating a bokeh effect.

arm

Scheduling

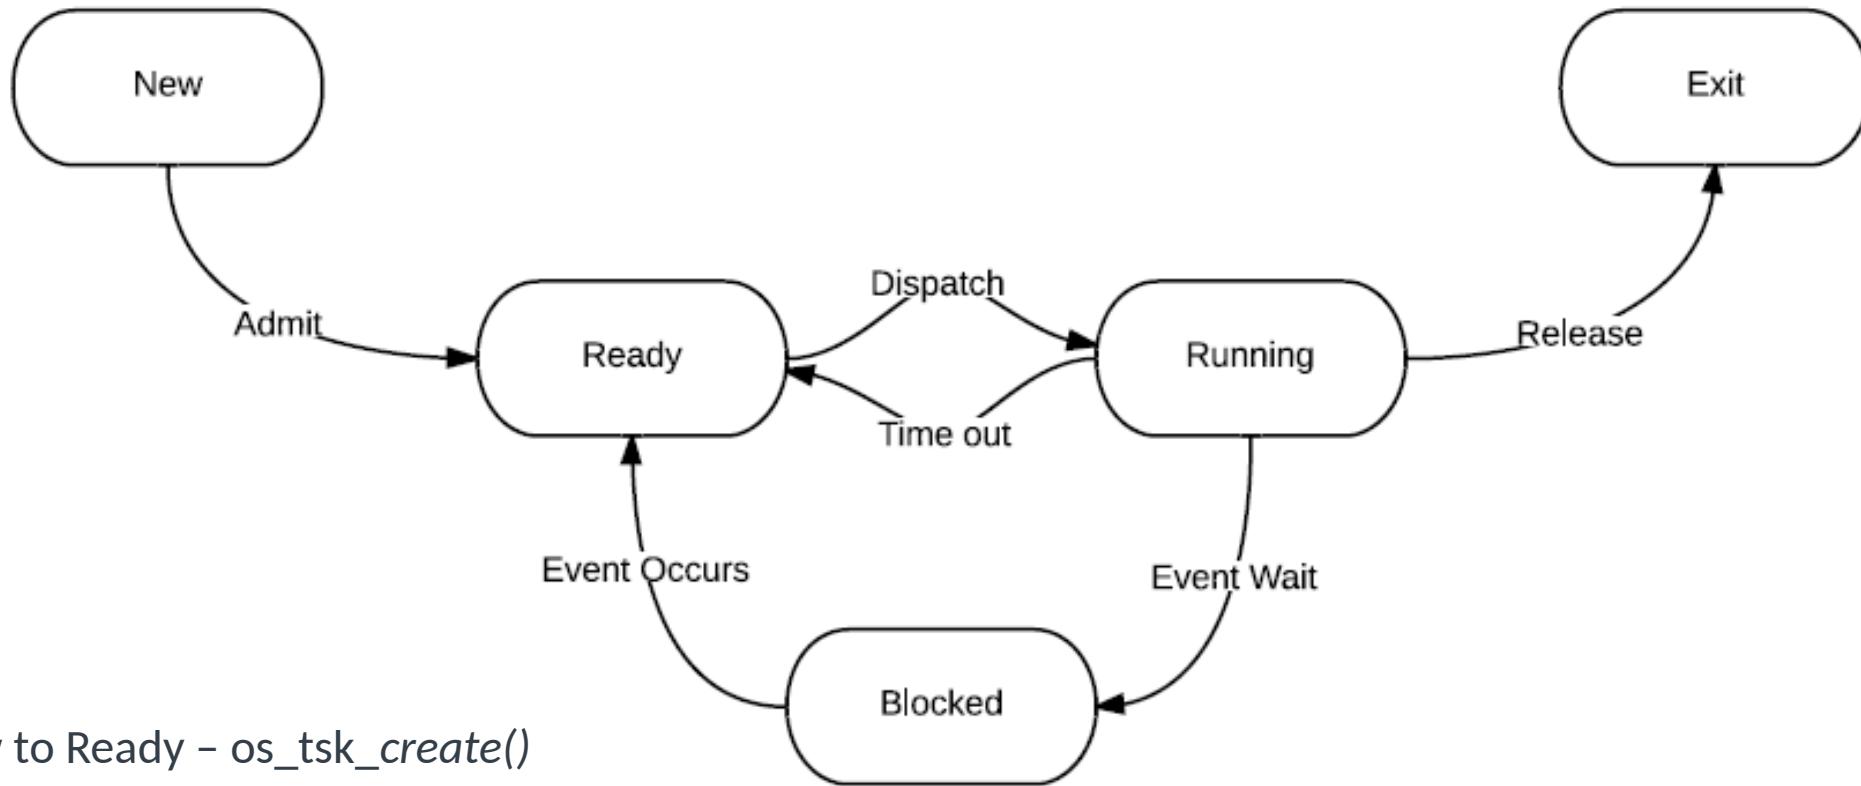
Previous

- Process
- Lab
 - Getting started with hardware using C and assembly
 - Context switching at assembly level
 - Basic RTX task operations

Current Module

- Levels of Scheduling
- Scheduling Algorithms
 - Non-pre-emptive
 - Pre-emptive
- Quantum Size
- Priority
- How to evaluate Performance
- Real Time Scheduling
- Aperiodic RT scheduling

Process State Model



- New to Ready - `os_tsk_create()`
- Running to Exit - `os_tsk_delete()`
- Ready to Running - `rt_dispatch()`
- Running to Ready - `os_tsk_pass()`
- When and How - the job of the scheduler
- To be practiced in lab session

Levels of Scheduling

- Long-term
 - Which processes to add to the ready queue
 - Historically the job of the operator
 - This scheduling level is unfrequently invoked compared to others
 - Establishes the degree of concurrency
 - Large scale computing systems
- Medium-term
 - Available for OSs that support swapping
 - Decides which process to swap out or suspend.
 - Moves tasks from the main memory to the secondary memory or vice versa
 - Uses the seven states model
- Short-term
 - Which ready process(es) to dispatch next
 - Performs context switching
- In this set of material, we will focus on the short-term scheduling. Whenever the term scheduling is used, unless otherwise specified, it refers to short-term scheduling.

What Makes a Good Scheduling Algorithm

- Scheduling criteria
 - Utilisation: computing resources, I/O, network
 - Throughput: completions of processes per time unit
 - Turnaround: the time it takes from New to Exit
 - Response: less delay in the interaction with users
 - Deadline: soft or hard
 - Priority: some processes are more important than the others
 - Fairness: despite the existence of privilege
 - Predictability: less variation in the performance
- A good scheduling algorithm is also application-specific

Behaviour Patterns of Processes

- Processes can be
 - CPU bound - more CPU computations than I/O interactions
 - IO bound - more IO interactions than CPU computations
- Dependency between processes
- Understanding the nature of the processes helps designing an optimized scheduling algorithm

Non-Preemptive Algorithms

Non-preemptive: a process keeps running until it voluntarily gives up or blocks.

Let's consider the single-CPU case

- **First-come, first-served (FCFS)**
 - Also known as first-in-first-out (FIFO)
 - Used for long-term scheduling by old batch systems
 - The ready queue needs to keep track of the task order.
 - `p_first = p_CB->p_lnk;`
 - `p_CB->p_lnk = p_first->p_lnk;`
 - Not really suitable for interactive systems:
 - Poor responsiveness and fairness
 - Unpredictable
 - Tends to starve short I/O-bound processes
 - Convoy Effect
 - Could be used as a sub-component in the queue

Non-Preemptive Algorithms

- **Shortest job first (SJF)**
 - Also known as shortest process next (SPN) or shortest job next (SJN)
 - Dispatches process with shortest execution time
 - To counteract the Convoy Effect, favours short processes
 - Minimizes the average waiting time
 - What if short processes keep coming?
 - Starves long processes
 - Also, how do you predict the execution time of the process?
 - Usually you can't, you can estimate
 - Exponentially weighted averaging

Exponentially Weighted Averaging

- You need to estimate the execution time of the process for some algorithms
- An iterative approach:
 - t_n : measure the n^{th} execution time
 - e_n : the estimate n^{th} execution time
 - w : the exponential weight ($0 < w < 1$)
- $e_{n+1} = wt_n + (1-w)e_n$
- $e_{n+1} = wt_n + (1-w)wt_{n-1} + (1-w)^2wt_{n-2} + \dots + (1-w)^iwt_{n-i} + \dots + (1-w)^{n+1}e_0$
- To simplify the implementation, take only the most recent n history executions into account
- Causes unpredictability

Preemptive Algorithms

Preemptive: processes can be interrupted by either processes with a higher priority or the task scheduler

- **Round-robin (RR)**
 - A preemptive FCFS variant
 - Needs to keep track of time
 - Once a process runs out of its time slice (or quantum), put it to the end of the Ready queue and dispatch the head of the queue
 - The Ready queue: a circular FCFS queue
 - Starvation-free, time-sharing
 - However, the average waiting time is increased, also frequent context switching results in higher overheads

Quantum Size

- How to pick the quantum size?
 - Large quantum size may starve small processes (degenerate to the FCFS)
 - Small quantum size introduces heavy context switching overhead
- Choose the size based on the needs of interaction
 - How frequently does the user interact with the OS?
 - The quantum size should provide the OS with good responsiveness to the user input
- Choose the size based on the performance (throughput)
 - Theoretical performance modelling, simulation
 - Practical experience
 - A rule of thumb: 80% of CPU bursts should be smaller than the quantum*.

Preemptive Algorithms

- **Shortest Remaining Time First (SRT)**
 - A preemptive SJF variant
 - Dispatch the process with shortest execution time once the current running process is finished or if any new task with shorter execution time is ready
 - You need to estimate the execution time as well

How does priority help?

- Is the shortest process always the most important?
 - On the contrary, usually trivial one
 - The user sets the importance of a process by assigning a priority to it
 - Can also use various types of scheduling algorithms
 - Static priority (remains unchanged throughout the lifecycle), dynamic priority(can be adjusted at runtime)
- Examples:
 - Multiple queues with different priorities, within the queue you may use FCFS or RR
 - In RR, higher priority processes may have larger quantum

Preemptive Algorithms

- Feedback: dynamic priorities
 - Most modern OSs use a variant of such algorithms
 - For multiple queues (presumably with different priorities), user can configure the following:
 - Number of queues
 - Which queue the process belongs to
 - When a process can be upgraded to a queue with a higher priority or downgraded to a queue with a lower priority
 - Scheduling algorithms for each queue
 - Scheduling algorithms for choosing a queue, e.g. a priority round-robin, 3 time slices for the highest, 2 time slices for normal and 1 for the lowest queue each round
 - Flexible and dynamic

A Feedback Example

- Multiple RR queues
- New processes join the queue with highest priority
- Each time a process is preempted, it will be downgraded and join the end of the queue with the next lower priority (except the lowest queue)
- Scheduler always dispatches the first process in the highest queue that is not empty.
- Problems?
 - Starvation in lower queues.
 - Lower priority has longer quantum

A Practical Problem: How To Evaluate The Performance

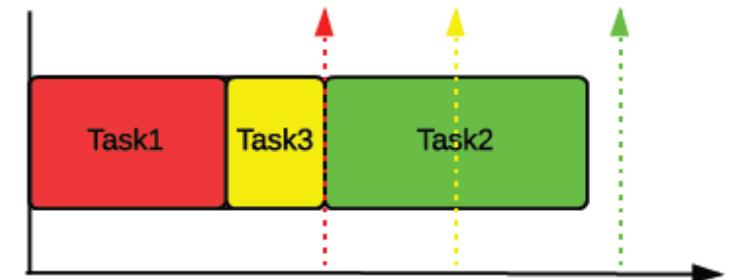
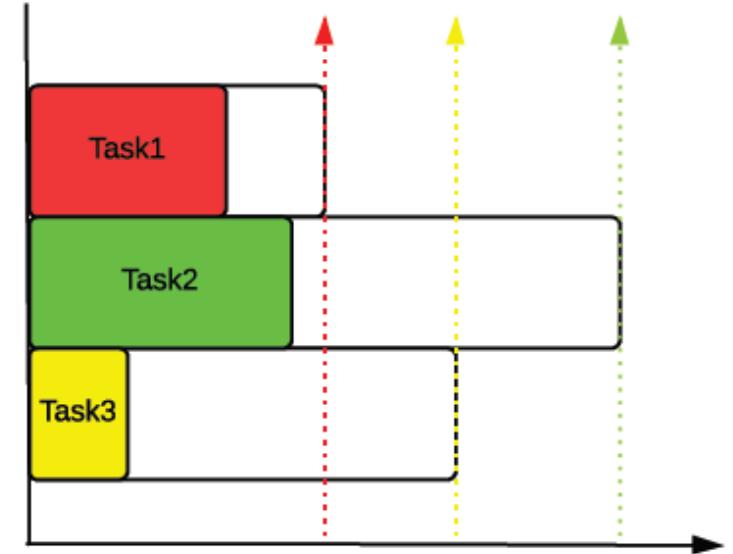
- Impossible to make definite comparisons in reality
- Either theoretical analysis or some empirical measurements
 - Theoretical analysis – Queuing network, stochastic petri net, process algebra or other modelling tools
 - Empirical measurements – Simulation

Real-Time Scheduling

- In a RTOS, meeting the deadline is more important than throughput and turn-around time
- Deadline: the maximal computational time allowed
- Needs special scheduling algorithms
- Preemption is usually a standard in RTOS: otherwise scheduler will have to be idle from time to time in case it misses tasks with earlier deadlines that arrive late.
- Two basic categories: aperiodic and periodic
- Also involves the priority and concurrency problems
- Here we look at simple aperiodic cases first
- All the plots in the following sections ignore the context switching time

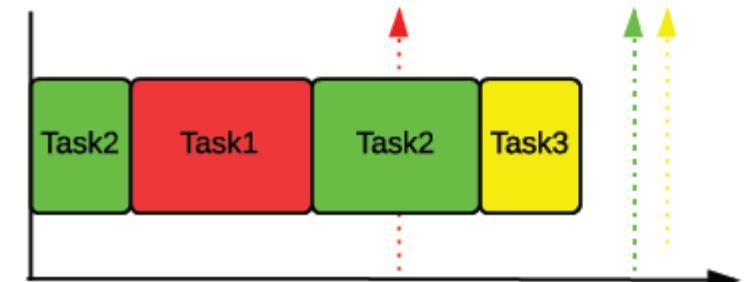
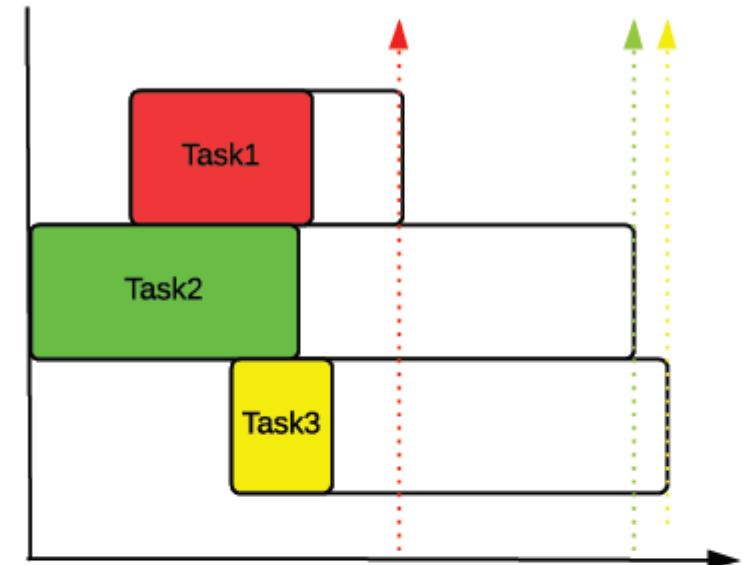
Aperiodic Real-Time Scheduling

- **Earliest Due Date (EDD)**
 - All tasks arrive at the same time
 - Known deadline and execution time for all tasks
 - Start the task with the earliest due date
 - Non-preemptive
 - Sorting algorithms
 - Optimal



Aperiodic Real-Time Scheduling

- **Earliest Deadline First (EDF)**
 - Different arrival times
 - Known deadline and execution time for all tasks
 - Start the task with earliest deadline
 - Preemptive (so overhead!)
 - Sorting algorithms
 - Optimal with respect to minimizing maximum lateness



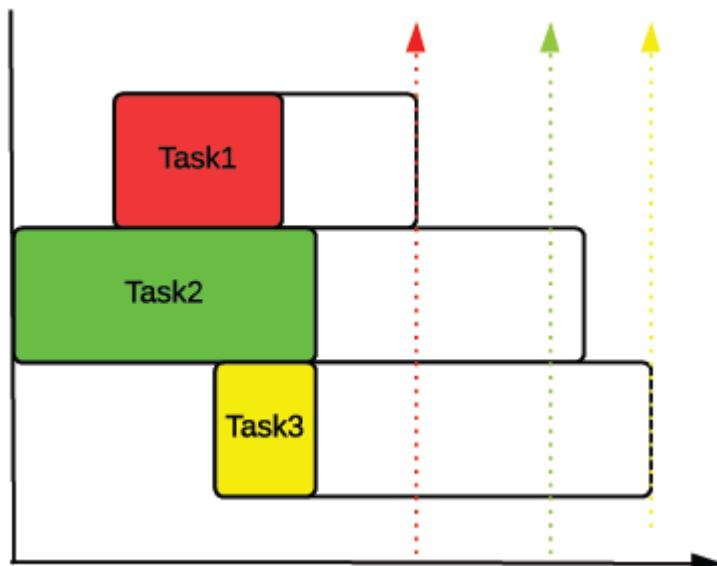
Aperiodic Real-Time Scheduling

- **Least Laxity (LL)**
 - Different arrival times
 - Laxity/slack: time to deadline minus (remaining) execution time
 - Known deadline and execution time for all tasks
 - Start the task with least laxity
 - Preemptive (more frequent)
 - Sorting algorithms
 - Dynamic priority: if not dispatched, then the priority will be raised; if it is executing, the priority maintains

Aperiodic Real-Time Scheduling

- An Example

Task	Starting at	Execution Time	Deadline
1	3	5	12
2	0	9	16
3	6	3	19

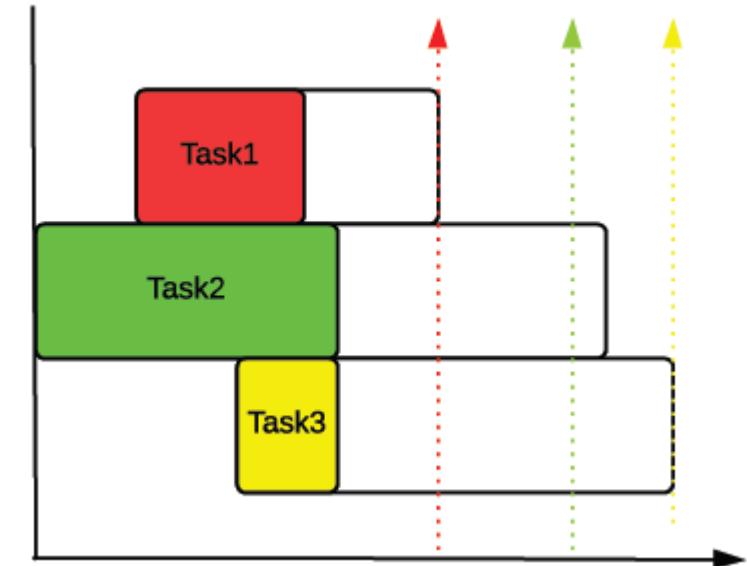


Aperiodic Real-Time Scheduling

- Try to fill in the laxity at each time step, dispatch the least

	0	1	2	3	4	5	6	7	8	9
1	-	-	-	4						
2	7	7	7	7						
3	-	-	-	-						

	10	11	12	13	14	15	16	17	18	19
1										
2										
3										



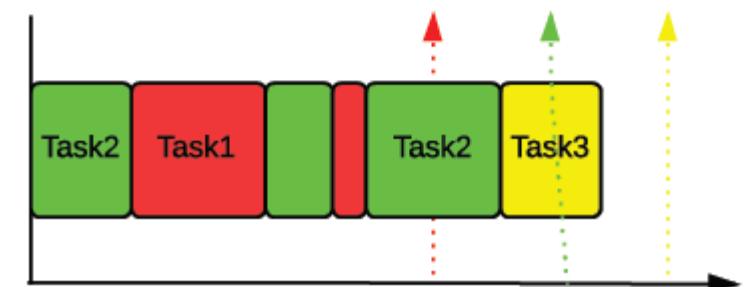
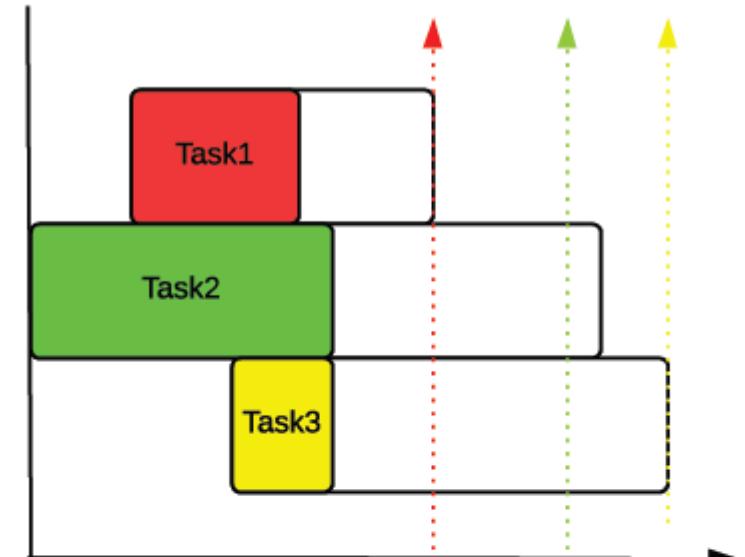
Aperiodic Real-Time Scheduling

- Underscore indicates the least laxity process

	0	1	2	3	4	5	6	7	8	9
1	-	-	-	4	4	4	4	4	3	2
2	7	7	7	7	6	5	4	3	3	3
3	-	-	-	-	-	-	10	9	8	7

	10	11	12	13	14	15	16	17	18	19
1	-	-	-	-	-	-	-	-	-	-
2	2	2	2	2	-	-	-	-	-	-
3	6	5	4	3	2	2	2	-	-	-

- Notice that at time step 7 and time step 9, preemptions happen due to change of priority (laxity)



All that we have discussed so far...

- All of the above is based on an ideal condition where:
 - Context switching time ignored
 - Single CPU: scheduling on multi-core is more challenging
 - “Well-behaved” processes
 - No dependency between processes
 - No periodic processes
- Be aware of the real world...

Next

- Assignment
 - Quantitative analysis and comparisons of various schedulers
- Concurrency