

A woman with dark hair and glasses, wearing a sequined top and pants, stands in a futuristic, glowing blue environment, possibly a server room or a digital space. She is looking down at her smartphone. The background is filled with vertical light streaks and glowing panels.

arm

RTX Task and Time Management

Content of this Module

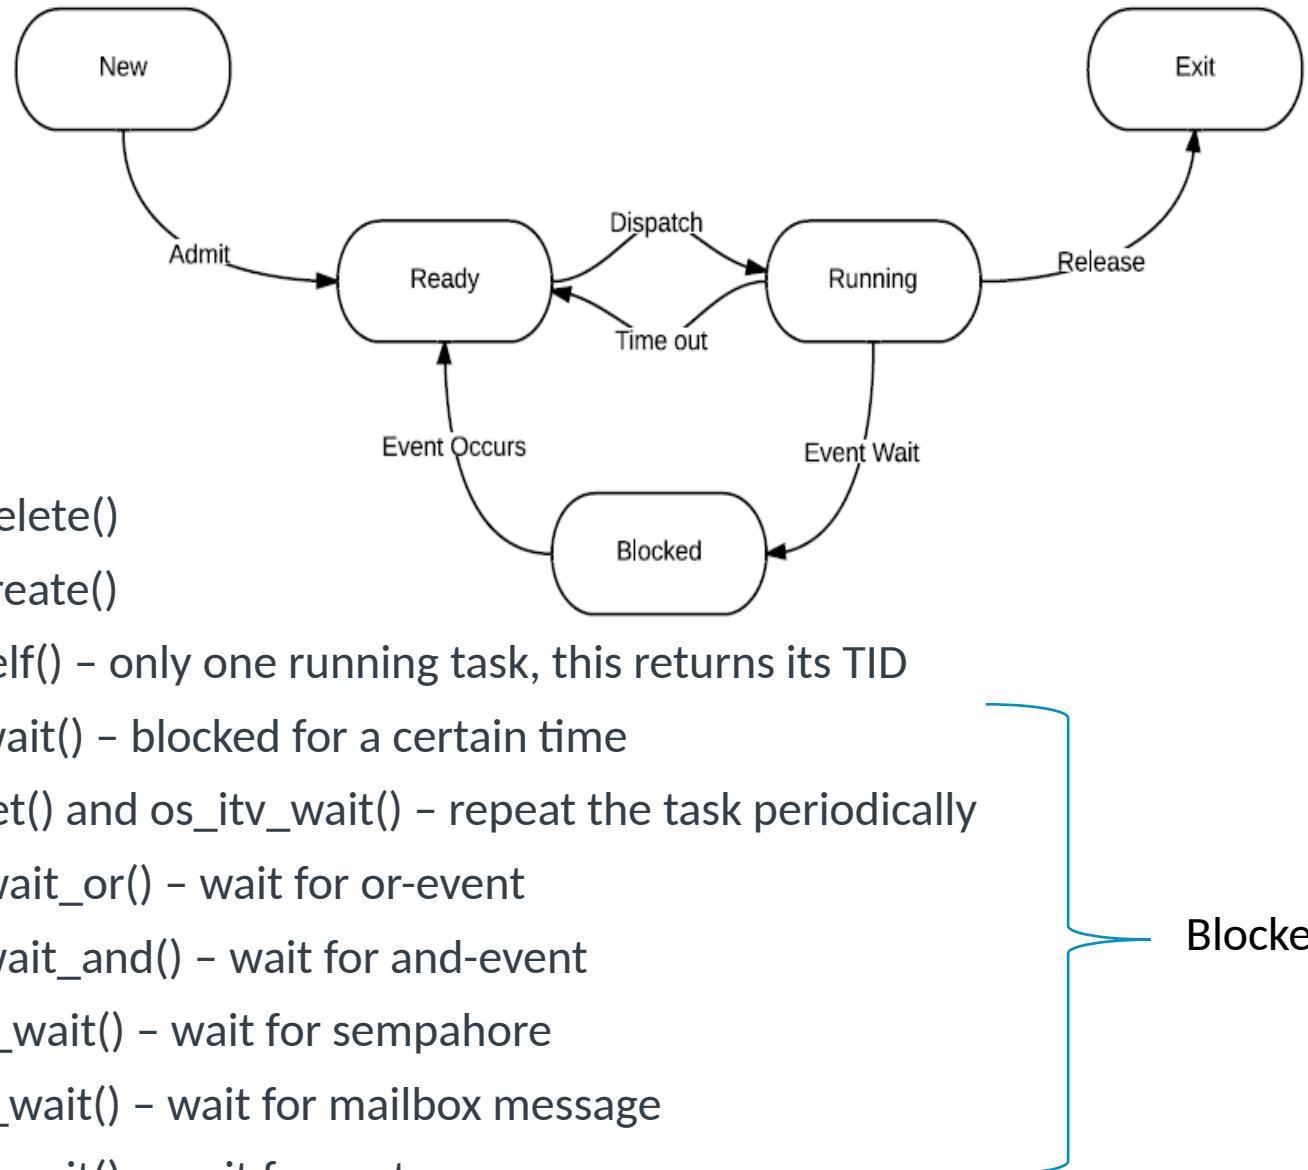
- Tasks in RTX, Task States, Creation and Deletion
- Advanced Real Time Task Scheduling
 - Periodic Task scheduling
 - Earliest Dead Line First
 - Rate Monotonic
 - System performance and Assumptions
 - Prioritized Scheduling
 - Priority Inversion
 - Priority Inheritance
 - Dead Lock
 - Priority Ceiling
- Priority Scheme in RTX, Options
- Simple Time Management API
- RTX control functions

Tasks in RTX

- Task is the term used to refer to a process in RTX
- C function with keyword `__task` preceding the function return type
- Some basic API functions:
 - `OS_TID`
 - `os_tsk_create(task, priority);`
 - `os_tsk_create_user(task, priority, &stack, sizeof(stack));`
 - `os_tsk_create_ex(task, priority, parameter);`
 - `os_sys_init(task);`
 - `os_tsk_delete(taskID);`
 - `os_tsk_delete_self();`

Task State

- P_TCB->state
 - rt_TypeDef.h
- #define INACTIVE 0; - Related function: os_tsk_delete()
- #define READY 1; - Related function: os_tsk_create()
- #define RUNNING 2; - Related function: os_tsk_self() – only one running task, this returns its TID
- #define WAIT_DLY 3; - Related function: os_dly_wait() – blocked for a certain time
- #define WAIT_ITV 4; - Related function: os_itv_set() and os_itv_wait() – repeat the task periodically
- #define WAIT_OR 5; - Related function: os_evt_wait_or() – wait for or-event
- #define WAIT_AND 6; - Related function: os_evt_wait_and() – wait for and-event
- #define WAIT_SEM 7; - Related function: os_sem_wait() – wait for semaphore
- #define WAIT_MBX 8; - Related function: os_mbx_wait() – wait for mailbox message
- #define WAIT_MUT 9; - Related function: os_mut_wait() – wait for mutex
 - rt_Task.h



Task Creation and Deletion

- Related functions

Task creation:

```
OS_TID - taskID;  
__task void some_task() {...}  
taskID = os_tsk_create (some_task, priority);  
taskID =  
os_tsk_create_user_ex(some_task,priority,&stack,sizeof(stack),parameters)
```

Task deletion:

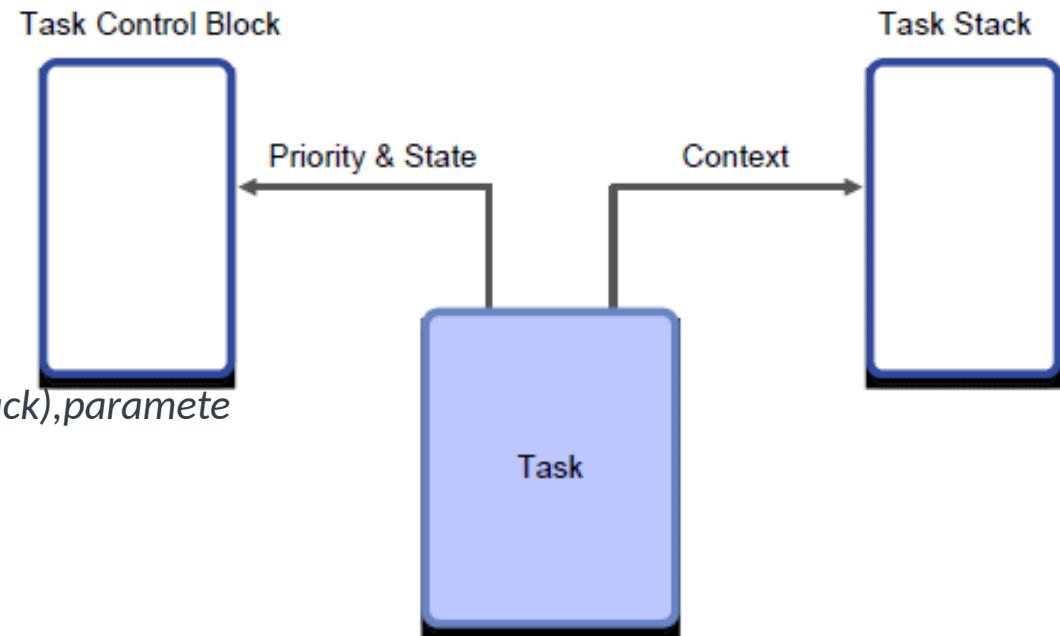
```
os_tsk_delete (some_task);
```

- Priority

- U8 member in the TCB in rt_TypeDef.h
- Larger number = higher priority
- '0' for an ideal task, '1' for a round-robin task

- User Stack (_user)

- You can set the default size in the configuration wizard



Task Creation and Deletion

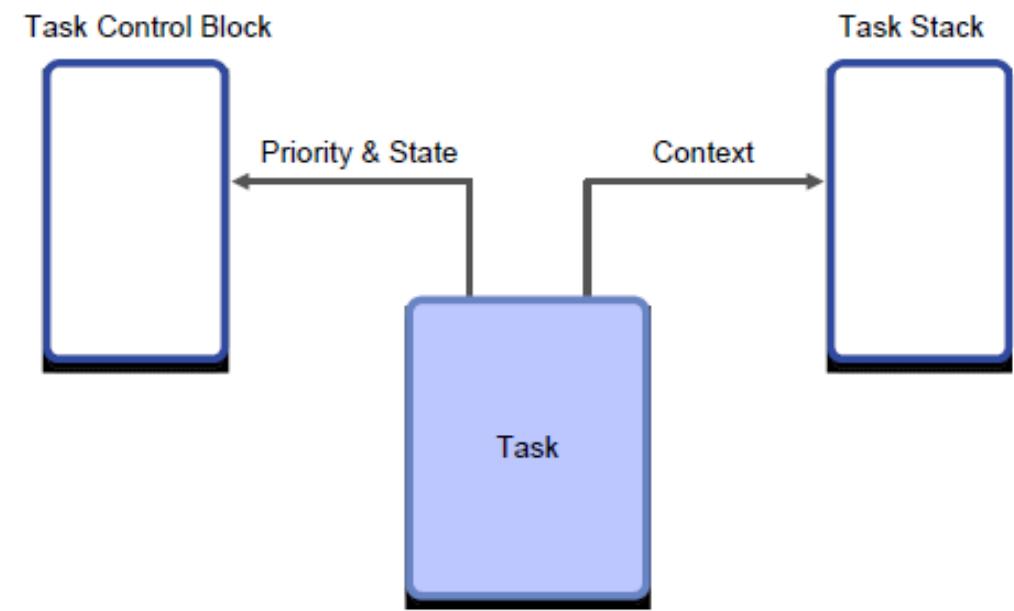
- Parameters (_ex)
 - RTX allows you to have multiple instances of the same task
 - Passing parameters to a task at creation, for example, to specify GPIO pins

```
taskID = os_tsk_create_ex(led_on_off, 0, led1)
taskID = os_tsk_create_ex(led_on_off, 0, led2)
```
- Tasks can be created and deleted by tasks
 - Tasks can delete other tasks and themselves

```
os_tsk_delete(some_task);
```
 - Tasks can create other tasks (but cannot create themselves)
 - The first task is created by:

```
os_sys_init(first_task);
```

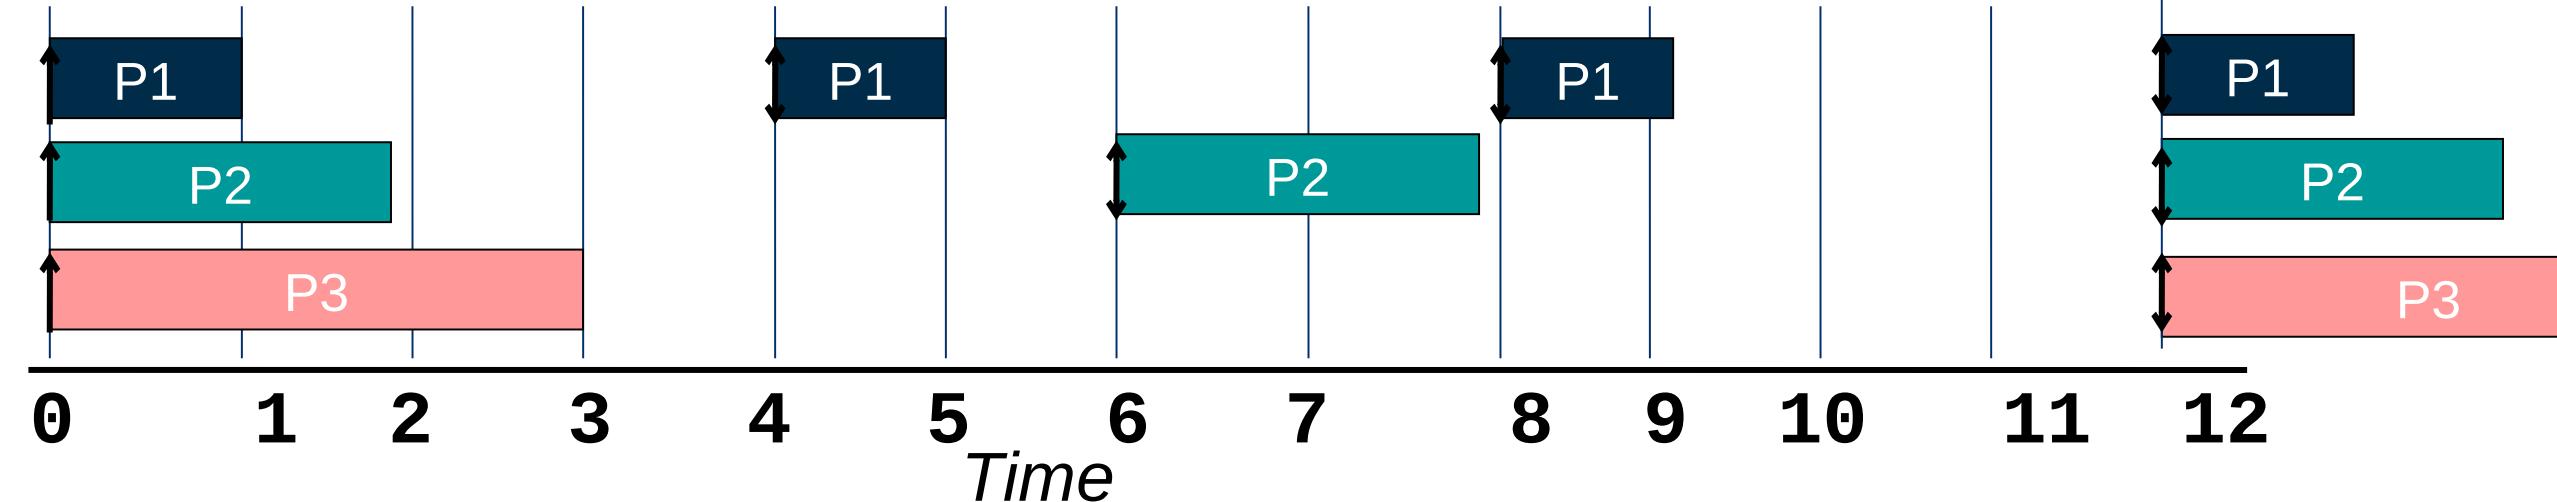
 - The first task is commonly used for
 - Initialization of hardware
 - Creation of all other tasks
 - Then deletes itself



Advanced Real-time Scheduling

- Recall the aperiodic real-time scheduling policies introduced previously
 - EDD, EDF and LL
 - Same priority
 - No inter-task dependency
 - Aperiodic
- Before introducing the RTX scheduler, we will discuss two more real-time scheduling scenarios
 - Periodic
 - Prioritized

Periodic Task Model



- Describes real-time computational requirements
- Constant process execution time C_i
- Tasks are periodic with period T_i
- Release time
 - Task is released every T_i time units and is ready to run
- Deadline D_i .
 - Might be related to period T_i making analysis easier
 - Can also be earlier than the end of a period!

Task	Exec. Time C_i	Period T_i	Deadline D_i
P1	1	4	4
P2	2	6	6
P3	3	12	12

Periodic Scheduling

- Goals
 - Meet all task deadlines!
 - Not always **schedulable** to meet all deadlines!
 - Maximize processor utilization (U)
 - The fraction of time CPU performs useful work
 - Limit scheduling overhead
 - Limit context switching overhead
- Which ready task to run?
 - Based only on task importance?
 - Would starve consistently low priority task unnecessarily
- Some common approaches
 - Earliest deadline first (EDF)
 - Rate monotonic (RM)
 - Deadline monotonic (DM) ($D_i \leq T_i$)

$$U = \sum_{i=1}^m \frac{C_i}{T_i}$$

Earliest Deadline First

- Run the job with the earliest deadline first
- Utilization-based schedulability test depends on deadline constraints
 - $D_i = T_i$: Schedulable if utilization $\leq 100\%$
 - $D_i > T_i$: Schedulable if utilization $\leq 100\%$
 - $D_i < T_i$: Have to use a more complicated test
- Hard to implement such dynamic policy
 - Run-time scheduler has to keep track of all deadlines, which also depends on tasks' release time
 - Sorting by deadline, time complexity: $O(n)$
- Periodic tasks only? Will repeat every hyper-period, the least common multiple of all task periods

Rate Monotonic

- Shortest period first, preemptive
- Optimal for workload where deadline = the end of period, i.e., $D_i=T_i$.
- Assume $D_i=T_i$, the utilization bound for RM: U_{max} is*

$$U_{Max} = n(2^{1/n} - 1)$$

- U_{max} approaches $\ln 2$ (around 0.693) as n approaches positive infinity.
- If $U < U_{max}$, 100% schedulable, otherwise, maybe schedulable. An example:

Task	Execution Time C	Period T	RM Dispatch
P1	1	8	First
P2	1	9	Second
P3	9	12	Third

- A modified approach: harmonic RM – task period can be evenly divide every longer task period for example, periods of 10, 20, 40, 80... Bound is 100%

¹¹ © 2021 ARM

*Originally proved by CL Liu and JW Layland in 1973. "Scheduling algorithms for multiprogramming in a hard-real-time environment."

System Performance During Transient Overload

- RM, DM – Each task has fixed priority. *So?*
 - This priority determines that tasks will be scheduled consistently
 - Task A will always preempt task B if needed
 - Task B will be forced to miss its deadline to help task A meet its deadline
- EDF – Each task has varying priority. *So?*
 - This priority depends upon when the task's deadline is, and hence when the task becomes ready to run (*release time*)
 - Task B may have higher priority than A depending on release times
 - To determine whether task A or B will miss its deadline we need to know their release times
 - Can theoretically harness up to 100% of the CPU performance

Our Assumptions

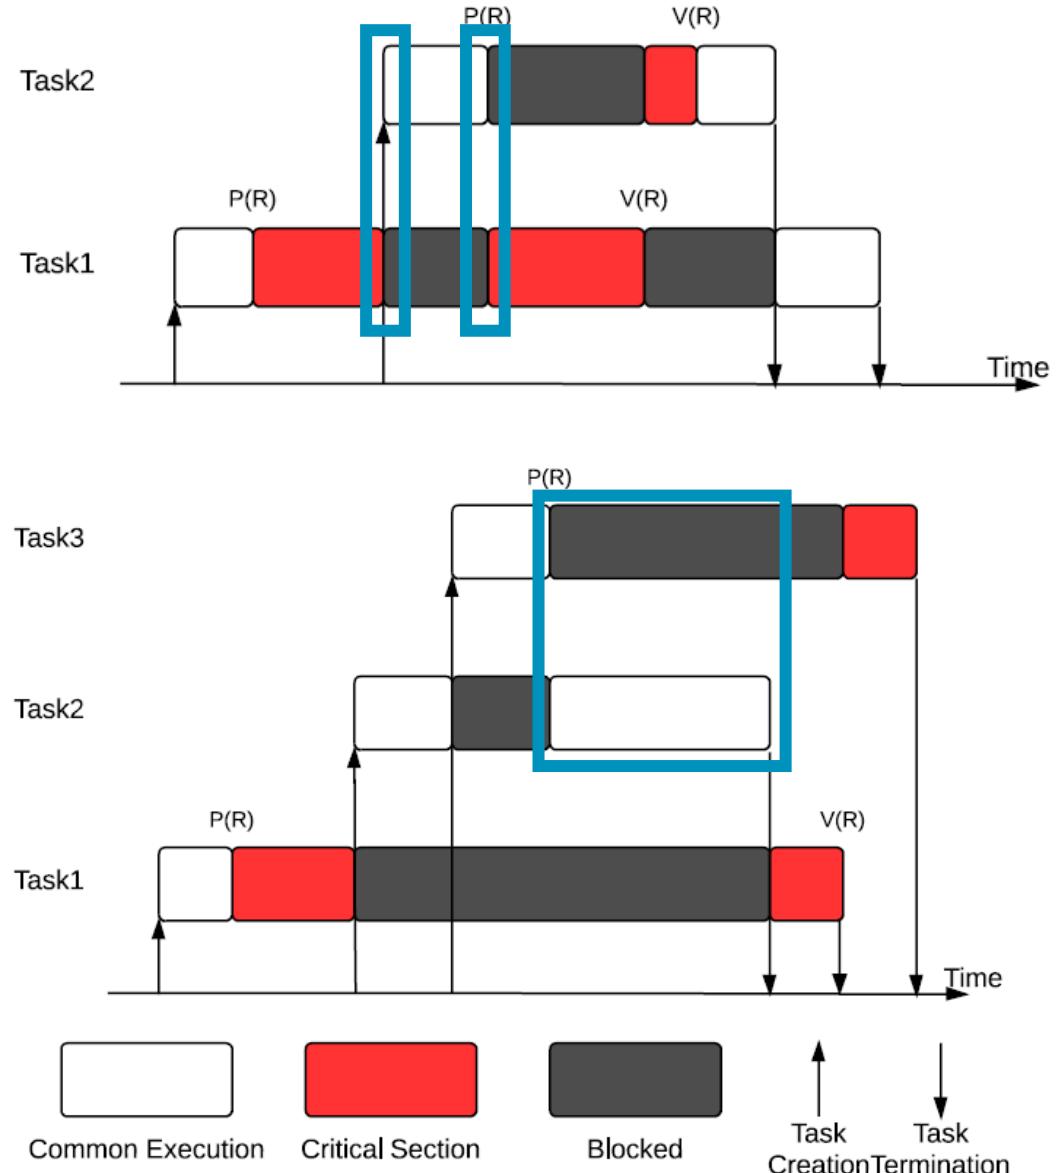
- Tasks are periodic with period T_i
 - Extension 1: Support aperiodic tasks by setting $T_i = \min(T_{\text{Interarrival}_i})$
- Single CPU
 - OK for some systems, but not all
- $C_{\text{ContextSwitch}} = 0$
 - Scheduler speed. EDF requires ordered list of ready tasks.
 - Task state information needs to be swapped
- Constant process execution time C_i
 - What if the worst-case time is much slower than average?
 - Then need a processor fast enough to handle the worst case
 - Most of the time the processor will be wasting resources
- No data dependencies between tasks
 - Not all tasks are ready to run simultaneously
 - Some tasks block on other tasks

Prioritized Scheduling

- `taskID = os_tsk_create(some_task, priority);`
- Prioritized Scheduling means that task with higher priority should be dispatched and can preempt lower priority tasks
- The problem is dependency. For instance, can a task in its critical section be preempted?
 - If no, any problem?
 - If yes, any problem?
 - The point of prioritized scheduling is based on the emergency level of the task, instead of whether it's in a critical section. But you have to pay extra attention to critical sections as otherwise they may cause conflict between tasks.
- **Priority inversion**

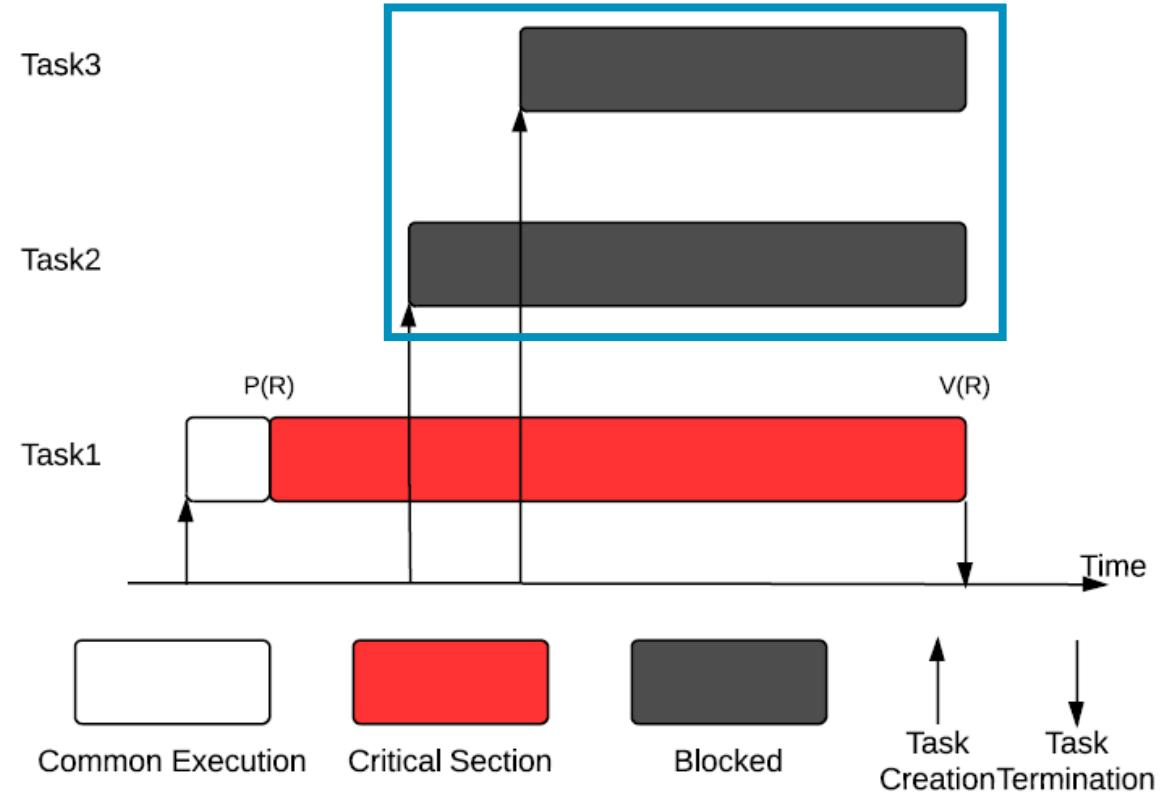
Preemptive Critical Section

- $\text{Priority}(T_N) > \text{Priority}(T_M)$, if $N > M$
- Seems okay?
- But what if a higher priority task would like to enter the same critical section? $P(R)$ blocks task 2 if task 1 has the mutex (recall the requirements for mutex). Still acceptable as contaminating CS may cause even worse disaster, better avoid that and let task 1 finish first.
- Blocking time seems to be bounded by the maximum length of critical section of lower priority tasks. But no. See the second example.
- Priority inversion: even though task 2 is not requesting R, it blocks task 3!



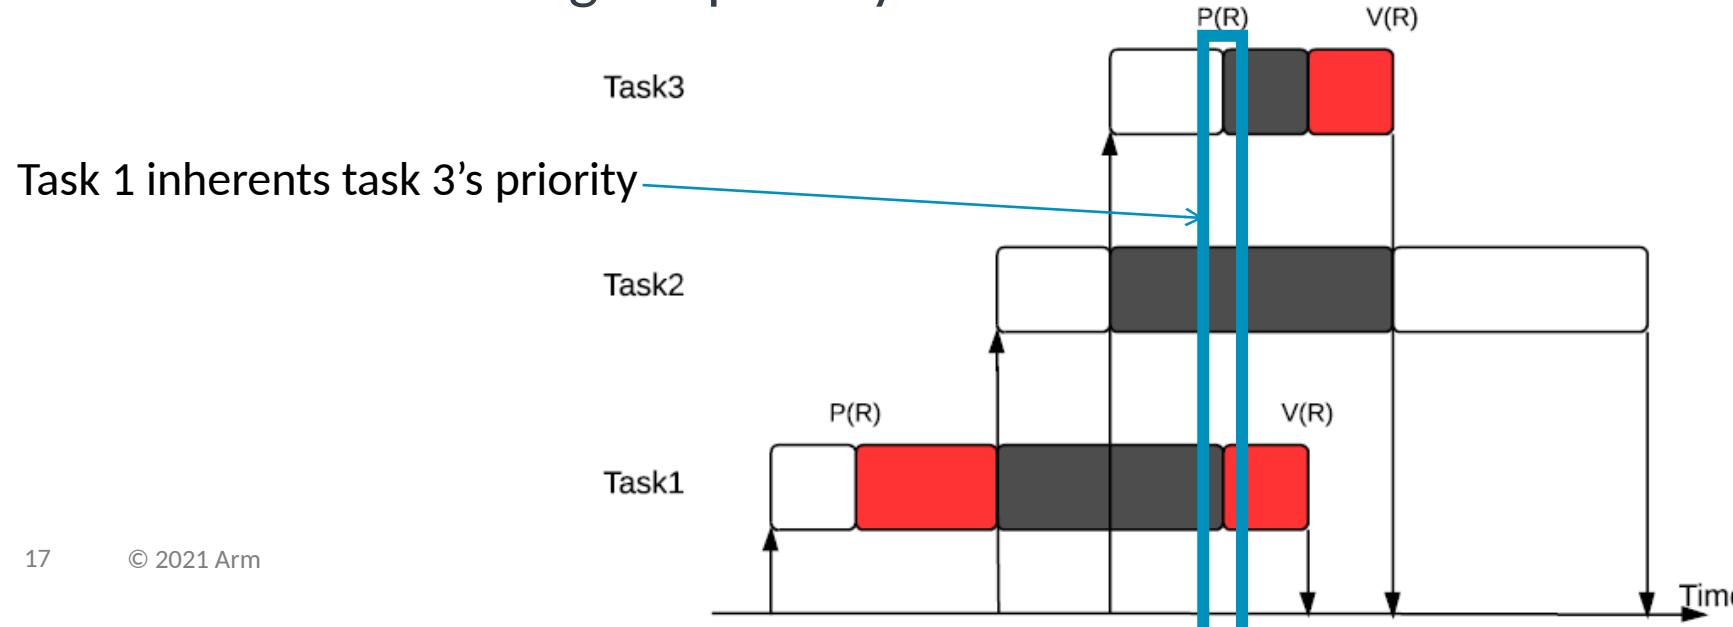
Non-Preemptive Critical Section

- Plausible? Can solve previous problem. But...
- Task 3 and 2 are on fire!



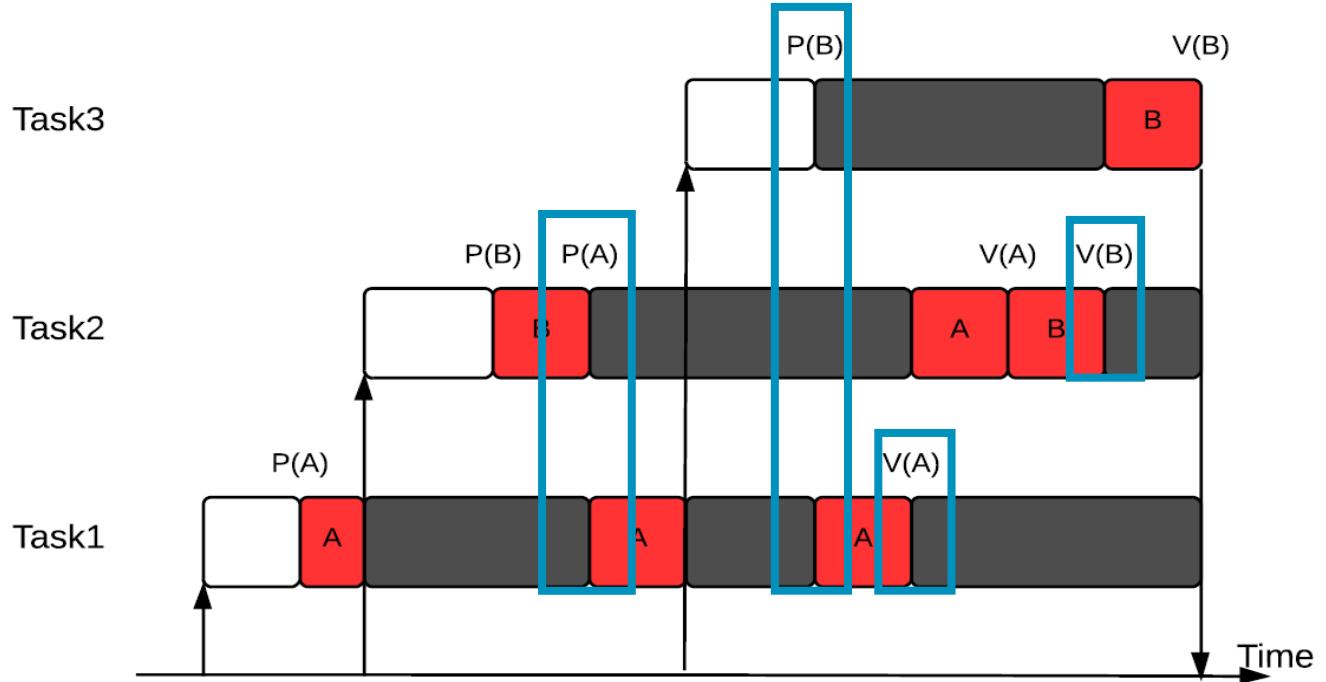
Priority Inheritance

- Applied in many RTOS including RTX.
- Blocking time is now bounded by the maximum (sum) length of critical section of lower priority tasks.
- The idea is to elevate the priority of low priority task (if it blocks high priority task) to the highest priority of tasks blocked by it.
- And resume its original priority when it exits the critical section.



Priority Inheritance

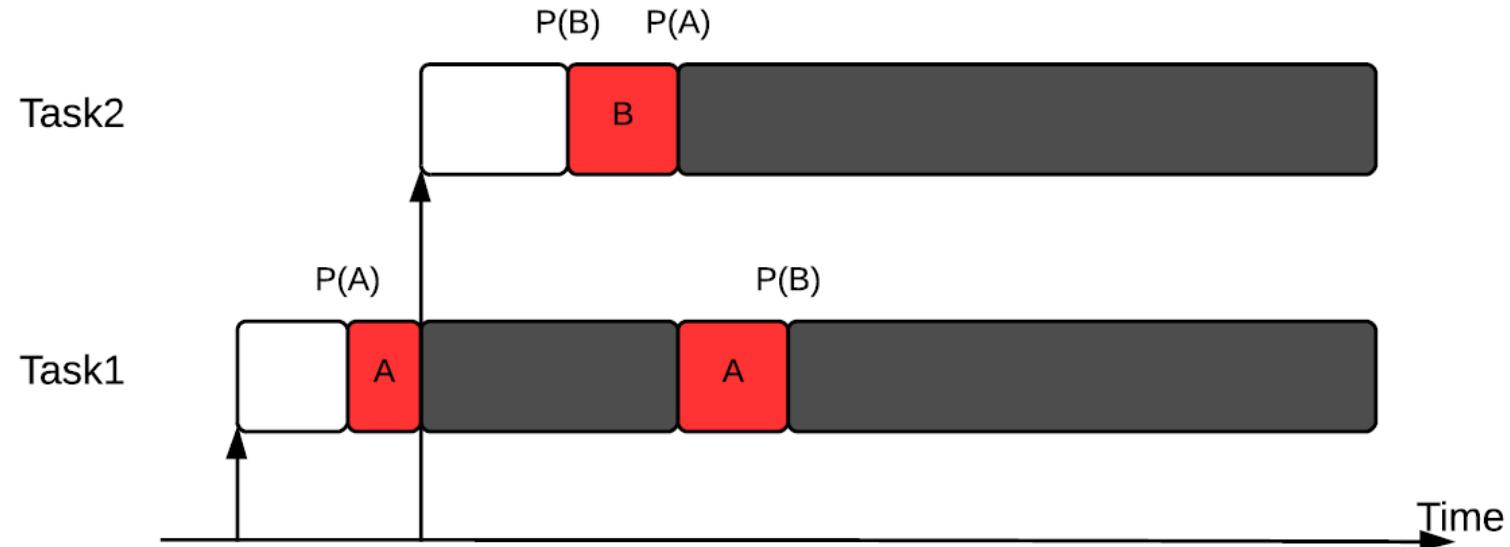
- If same priority, then FIFO
- Another example:



- Task 2 blocked by task 1, hence $\text{Priority}(T_1) = \text{Priority}(T_2)$
- Then, task 3 is blocked by task 1, hence $\text{Priority}(T_1) = \text{Priority}(T_3)$, also $\text{Priority}(T_1)=\text{Priority}(T_2)=\text{Priority}(T_3)$ as task 2 is blocked by task 1. So Task 1 continues(FIFO)
- After task 1 releases A, its priority will be resumed. Task 2 continues.
- Similarly, after task 2 releases B, its priority will be downgraded and task 3 continues.

Deadlock Caused by Priority Inheritance

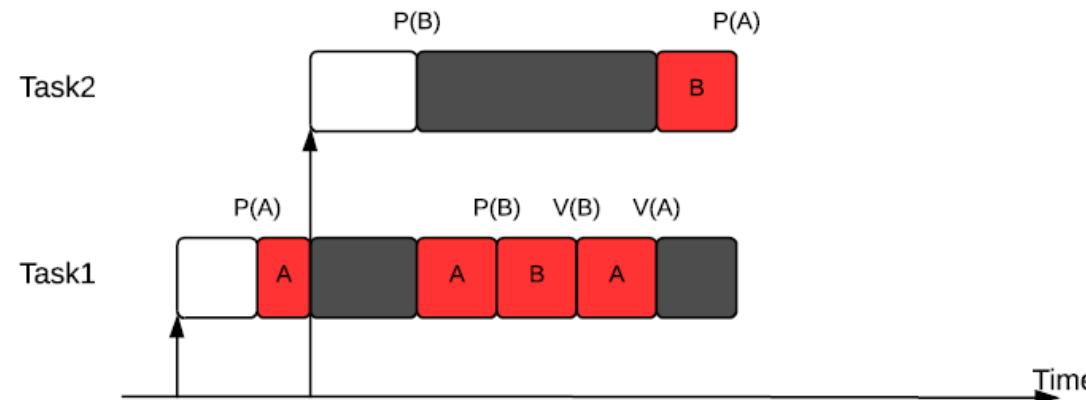
- Unfortunately, priority inheritance is not perfect. Consider the following example:



- Recall Coffman conditions for deadlock. Deadlock here!

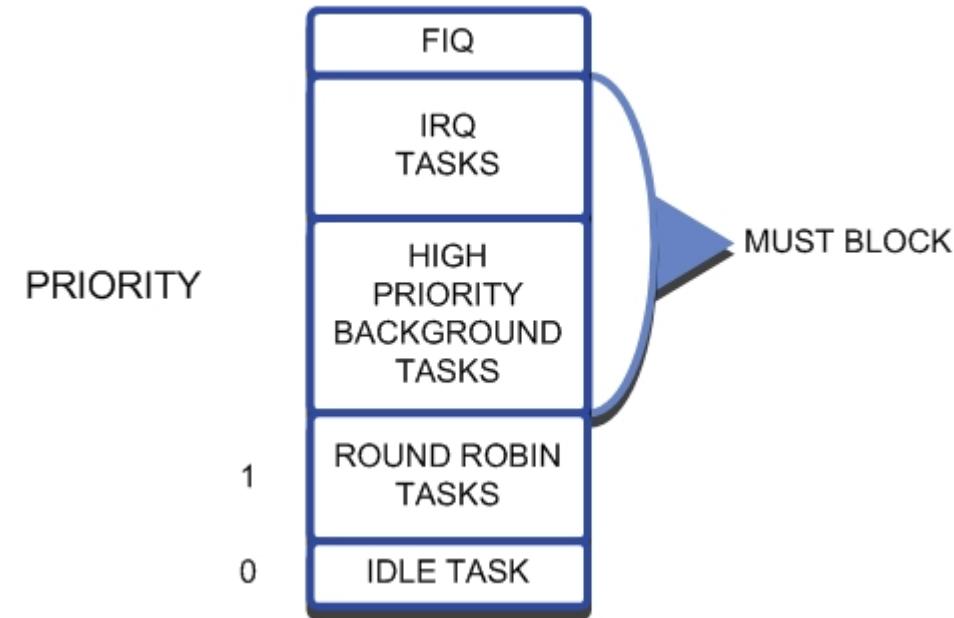
Priority Ceiling

- Improvement on priority inheritance. But **not implemented in the RTX**.
- Define the priority ceiling for a resource: $PC(R)$ are the highest priority processes that may request mutual exclusive access to R
- Priority ceiling accepts requests to enter any critical section only if task A's priority is **larger** than PC (all resources accessed by tasks other than A). This solves deadlocks by preventing circular wait conditions
- If task A is blocked due to this, then the priority of task B which is blocking task A, will be elevated to A's priority (if A is the highest) as the usual priority inheritance protocol.
- Another improvement is that it reduces the blocking time to at most the length of one critical section of a low priority task.



Task Priority Scheme in RTX

- RTX supports priority scheduling and applies priority inheritance policy with mutex.
- Possible to alter the priority of a task after creation.
`OS_RESULT os_tsk_prio(taskID, priority);`
`OS_RESULT os_tsk_prio_self(priority);`
- Higher number, higher priority, from 0 to 255
 - A FIFO queue for each priority
 - Fast interrupt has the highest priority (handled by CM core instead of RTX)
 - Then interrupts
 - Then tasks with a priority larger than 1. RTX will not preempt tasks with priority larger than 1 so they must be self-blocking!
 - Tasks in queue 0 will be put into queue 1 when there are no tasks running



RTX Scheduling Options

- Pre-emptive scheduling
 - Each task has a **different priority** and will run until it is pre-empted or has reached a blocking OS call.
- Round-Robin scheduling
 - Each task has the **same priority** and will run for a fixed period, or time slice, or until it has reached a blocking OS call.
 - Quantum is determined by the **RTX_Conf_CM.c** file: Round-Robin Timeout [ticks] × Timer tick value.
For example, by default, the time quantum is $10000\mu\text{s} \times 5 = 50000\mu\text{s}=50\text{ms}$.
 - If quantum expired, state will be changed to READY.
- Co-operative multi-tasking
 - Each task has the **same priority** and the Round-Robin is disabled. Each task will run until it reached a blocking OS call or uses the `os_tsk_pass()` call.
- The default scheduling option for RTX is **Round-Robin Pre-emptive**. Prioritized RR.
- RTX does not provide aperiodic or periodic deadline scheduling.

Simple Time Management APIs

- Besides the scheduling options, RTX can also manage time through the following APIs:
- Get current time
 - `U32 os_time_get (void);`
 - Returns current OS time measured in ticks
- Simple delay
 - `void os_dly_wait (unsigned short delay_time)`
 - Turn the state of calling task into WAIT_DLY
 - Actual delay time = $delay_time \times$ timer tick value
 - After the delay time, turn the state of the task back to READY
 - Make sure you have the correct timer clock value
- Periodic task execution
 - Set the interval within the task by calling `void os_itv_set (unsigned short interval_time)` first
 - Actual delay time = $interval_time \times$ timer tick value
 - Then change the to WAIT_ITV by calling `void os_itv_wait (void)` inside the task

Simple Time Management APIs

- Cooperate function
 - `void os_tsk_pass();`
 - Task calling this function will voluntarily give up its current possession of CPU and allow next ready task with same priority to run.
 - No task switch if the next ready task has a lower priority than the calling task.
- User timer function
 - Set the number of available user timer first.
 - Set a single shot timer by calling `OS_ID os_tmr_create (unsigned short tcnt, unsigned short info)`
 - Actual countdown time = $tcnt \times$ timer tick value
 - After it reaches zero, RTX runs the callback function `void os_tmr_call (U16 info)`. The prototype function can be found in `RTX_Conf_CM.c`
 - *Info* is the parameter passed to the callback function

Other RTX Control Functions

- Scheduler lock – stops task switching, though ISRs can still execute
 - If you want your code to run as a contiguous block, (critical section for example), it is a good idea to temporally stop the scheduler.
 - `void os_tsk_lock();`
 - Disables kernel timer interrupts, stopping task switching
 - `void os_tsk_unlock();`
 - Enables kernel timer interrupts, allowing task switching
- Power management - used in low power RTX
 - Used with `os_dly_wait` or `os_tmr_create`. Hibernate the CPU until next activity.
 - `U32 os_suspend();`
 - Calculates for how many ticks the system can sleep
 - Locks task scheduler
 - Use this in `idle_task` to determine when to set wake-up timer before sleeping (`wfi` instruction)
 - `void os_resume(U32 sleep_time);`
 - Updates system timer with sleep duration (number of ticks)
 - Unlocks task scheduler, allowing normal scheduling
 - Use this in `idle_task` to resume

Next

- Lab
- Sharing data in RTX